

1. Pandas DataFrame Basics

1.1 Introduction

Pandas is an open source Python library for data analysis. It gives Python the ability to work with spreadsheet-like data for fast data loading, manipulating, aligning, and merging, among other functions. To give Python these enhanced features, Pandas introduces two new data types to Python: `Series` and `DataFrame`. The `DataFrame` represents your entire spreadsheet or rectangular data, whereas the `Series` is a single column of the `DataFrame`. A Pandas `DataFrame` can also be thought of as a dictionary or collection of `Series` objects.

Why should you use a programming language like Python and a tool like Pandas to work with data? It boils down to automation and reproducibility. If a particular set of analyses need to be performed on multiple data sets, a programming language has the ability to automate the analysis on those data sets. Although many spreadsheet programs have their own macro programming languages, many users do not use them. Furthermore, not all spreadsheet programs are available on all operating systems. Performing data analysis using a programming language forces the user to maintain a running record of all steps performed on the data. I, like many people, have accidentally hit a key while viewing data in a spreadsheet program, only to find out that my results no longer make any sense due to bad data. This is not to say that spreadsheet programs are bad or that they do not have their place in the data workflow, they do. Rather, my point is that there are better and more reliable tools out there.

Concept Map

1. Prior knowledge needed (appendix)

a. relative directories

b. calling functions

c. dot notation

d. primitive Python containers

e. variable assignment

2. This chapter

a. loading data

b. subset data

c. slicing

d. filtering

e. basic Pandas data structures (`Series` , `DataFrame`)

f. resemble other Python containers (`list` , `numpy.ndarray`)

g. basic indexing

Objectives

This chapter will cover:

1. Loading a simple delimited data file

2. Counting how many rows and columns were loaded

3. Determining which type of data was loaded

4. Looking at different parts of the data by subsetting rows and columns

1.2 Loading Your First Data Set

When given a data set, we first load it and begin looking at its structure and contents. The simplest way of looking at a data set is to examine and subset specific rows and columns. We can see which type of information is stored in each column, and can start looking for patterns by aggregating descriptive statistics.

Since Pandas is not part of the Python standard library, we have to first tell Python to load (`import`) the library.

```
import pandas
```

With the library loaded, we can use the `read_csv` function to load a CSV data file. To access the `read_csv` function from Pandas, we use dot notation. More on dot notations can be found in [Appendices H, O, and S](#).

ABOUT THE GAPMINDER DATA SET

The Gapminder data set originally comes from www.gapminder.org. The version of the Gapminder data used in this book was prepared by Jennifer Bryan from the University of British Columbia. The repository can be found at: [www.github.com/jennybc/gapminder](https://github.com/jennybc/gapminder).

[Click here to view code image](#)

```
# by default the read_csv function will read a comma-separated file;  
  
# our Gapminder data are separated by tabs  
  
# we can use the sep parameter and indicate a tab with \t  
  
df = pandas.read_csv('../data/gapminder.tsv', sep='\t')  
  
# we use the head method so Python shows us only the first 5 rows  
  
print(df.head())
```

	country	continent	year	lifeExp	pop	gdpPerCap
0	Afghanistan	Asia	1952	28.801	8425333	779.445314
1	Afghanistan	Asia	1957	30.332	9240934	820.853030
2	Afghanistan	Asia	1962	31.997	10267083	853.100710
3	Afghanistan	Asia	1967	34.020	11537966	836.197138
4	Afghanistan	Asia	1972	36.088	13079460	739.981106

When working with Pandas functions, it is common practice to give `pandas` the alias `pd`. Thus the following code is equivalent to the preceding example:

[Click here to view code image](#)

```
import pandas as pd

df = pd.read_csv('../data/gapminder.tsv', sep='\t')
```

We can check whether we are working with a Pandas `DataFrame` by using the built-in `type` function (i.e., it comes directly from Python, not any package such as Pandas).

[Click here to view code image](#)

```
print(type(df))

<class 'pandas.core.frame.DataFrame'>
```

The `type` function is handy when you begin working with many different types of Python objects and need to know which object you are currently working on.

The data set we loaded is currently saved as a Pandas `DataFrame` object and is relatively small. Every `DataFrame` object has a `shape` attribute that will give us the number of rows and columns of the `DataFrame`.

[Click here to view code image](#)

```
# get the number of rows and columns  
  
print(df.shape)
```

```
| (1704, 6)
```

The `shape` attribute returns a tuple ([Appendix J](#)) in which the first value is the number of rows and the second number is the number of columns. From the preceding results, we see our Gapminder data set has 1704 rows and 6 columns.

Since `shape` is an attribute of the dataframe, and not a function or method of the `DataFrame`, it does not have parentheses after the period. If you made the mistake of putting parentheses after the `shape` attribute, it would return an error.

[Click here to view code image](#)

```
# shape is an attribute, not a method  
  
# this will cause an error  
  
print(df.shape())
```

```
| Traceback (most recent call last):
```

```
|   File "<ipython-input-1-e05f133c2628>", line 2, in <module>
```

```
|     print(df.shape())
```

```
|   TypeError: 'tuple' object is not callable
```

Typically, when first looking at a data set, we want to know how many rows and columns there are (we just did that). To get the gist of which information it contains, we look at the columns. The column names, like `shape`, are specified using the `column` attribute of the `Dataframe` object.

[Click here to view code image](#)

```
# get column names
```

```
print(df.columns)
```

```
Index(['country', 'continent', 'year', 'lifeExp', 'pop',
       'gdpPercap'],
      dtype='object')
```

QUESTION

What is the `type` of the column names?

The Pandas `DataFrame` object is similar to the `DataFrame`-like objects found in other languages (e.g., Julia and R). Each column (`Series`) has to be the same type, whereas each row can contain mixed types. In our current example, we can expect the `country` column to be all strings and the `year` to be integers. However, it's best to make sure that is the case by using the `dtypes` attribute or the `info` method. [Table 1.1](#) compares the types in Pandas to the types in native Python.

[Click here to view code image](#)

```
# get the dtype of each column
```

```
print(df.dtypes)
```

```
country      object  
continent    object  
  
year         int64  
  
lifeExp     float64  
  
pop          int64  
  
gdpPercap   float64  
  
dtype: object
```

```
# get more information about our data
```

```
print(df.info())
```

```
<class 'pandas.core.frame.DataFrame'>  
  
RangeIndex: 1704 entries, 0 to 1703  
  
Data columns (total 6 columns):  
  
country      1704 non-null object  
  
continent    1704 non-null object  
  
year         1704 non-null int64  
  
lifeExp     1704 non-null float64  
  
pop          1704 non-null int64  
  
gdpPercap   1704 non-null float64  
  
dtypes: float64(2), int64(2), object(2)  
  
memory usage: 80.0+ KB
```

None

Table 1.1 Pandas Types Versus Python Types

Pandas Type	Python Type	Description
object	string	Most common data type
int64	int	Whole numbers
float64	float	Numbers with decimals
datetime64	date-time	<code>datetime</code> is found in the Python standard library (i.e., it is not loaded by default and needs to be imported)

1.3 Looking at Columns, Rows, and Cells

Now that we're able to load a simple data file, we want to be able to inspect its contents. We could `print` out the contents of the dataframe, but with today's data, there are often too many cells to make sense of all the printed information. Instead, the best way to look at our data is to inspect it in parts by looking at various subsets of the data. We already saw that we can use the `head` method of a dataframe to look at the first five rows of our data. This is useful to see if our data loaded properly and to get a sense of each of the columns, its name, and its contents. Sometimes, however, we may want to see only particular rows, columns, or values from our data.

Before continuing, make sure you are familiar with Python containers ([Appendices I, J, and K](#)).

1.3.1 Subsetting Columns

If we want to examine multiple columns, we can specify them by names, positions, or ranges.

1.3.1.1 Subsetting Columns by Name

If we want only a specific column from our data, we can access the data using square brackets.

[Click here to view code image](#)

```
# just get the country column and save it to its own variable  
  
country_df = df[ 'country' ]  
  
# show the first 5 observations  
  
print(country_df.head())  
  
0    Afghanistan  
1    Afghanistan  
2    Afghanistan  
3    Afghanistan  
4    Afghanistan  
  
Name: country, dtype: object
```

```
# show the last 5 observations
```

```
print(country_df.tail())
```

```
1699    Zimbabwe  
1700    Zimbabwe  
1701    Zimbabwe  
1702    Zimbabwe  
1703    Zimbabwe  
  
Name: country, dtype: object
```

To specify multiple columns by the column name, we need to pass in a Python `list` between the square brackets. This may look a bit strange since there will be two sets of square brackets.

[Click here to view code image](#)

```
# Looking at country, continent, and year  
  
subset = df[['country', 'continent', 'year']]  
  
print(subset.head())
```

```
country continent year  
0 Afghanistan Asia 1952  
1 Afghanistan Asia 1957  
2 Afghanistan Asia 1962  
3 Afghanistan Asia 1967  
4 Afghanistan Asia 1972
```

```
print(subset.tail())
```

	country	continent	year
1699	Zimbabwe	Africa	1987
1700	Zimbabwe	Africa	1992
1701	Zimbabwe	Africa	1997
1702	Zimbabwe	Africa	2002
1703	Zimbabwe	Africa	2007

Again, you can opt to `print` the entire `subset` dataframe. We won't use this option in this book, as it would take up an unnecessary amount of space.

1.3.1.2 Subsetting Columns by Index Position Break in Pandas v0.20

At times, you may want to get a particular column by its position, rather than its name. For example, you want to get the first ("country") column and third column ("year"), or just the last column ("gdpPercap").

As of pandas `v0.20`, you are no longer able to pass in a list of integers in the square brackets to subset columns. For example, `df[[1]]`, `df[[0, -1]]`, and `df[list(range(5))]` no longer work. There are other ways of subsetting columns ([Section 1.3.3](#)), but they build on the technique used to subset rows.

1.3.2 Subsetting Rows

Rows can be subset in multiple ways, by row name or row index. [Table 1.2](#) gives a quick overview of the various methods.

Table 1.2 Different Methods of Indexing Rows (or Columns)

Subset method	Description
loc	Subset based on index label (row name)
iloc	Subset based on row index (row number)
ix (no longer works in Pandas v0.20)	Subset based on index label or row index

1.3.2.1 Subset Rows by Index Label: loc

Let's take a look at part of our Gapminder data.

[Click here to view code image](#)

```
print(df.head())
```

	country	continent	year	lifeExp	pop	gdpPercap
0	Afghanistan	Asia	1952	28.801	8425333	779.445314
1	Afghanistan	Asia	1957	30.332	9240934	820.853030
2	Afghanistan	Asia	1962	31.997	10267083	853.100710
3	Afghanistan	Asia	1967	34.020	11537966	836.197138
4	Afghanistan	Asia	1972	36.088	13079460	739.981106

On the left side of the printed dataframe, we see what appear to be row numbers. This column-less row of values is the index label of the dataframe. Think of the index label as being like a column name, but for

rows instead of columns. By default, Pandas will fill in the index labels with the row numbers (note that it starts counting from 0). A common example where the row index labels are not the same as the row number is when we work with time series data. In that case, the index label will be a timestamp of sorts. For now, though, we will keep the default row number values.

We can use the `loc` attribute on the dataframe to subset rows based on the index label.

[Click here to view code image](#)

```
# get the first row
```

```
# Python counts from 0
```

```
print(df.loc[0])
```

```
country      Afghanistan
continent      Asia
year          1952
lifeExp        28.801
pop            8425333
gdpPercap      779.445
Name: 0, dtype: object
```

```
# get the 100th row
```

```
# Python counts from 0
```

```
print(df.loc[99])
```

```
country      Bangladesh
continent     Asia
year         1967
lifeExp      43.453
pop          62821884
gdpPercap    721.186
Name: 99, dtype: object
```

```
# get the last row
# this will cause an error
print(df.loc[-1])
```

```
Traceback (most recent call last):
```

```
  File "/home/dchen/anaconda3/envs/book36/lib/python3.6/site-
  packages/pandas/core/indexing.py", line 1434, in _has_valid_type
    error()
```

```
KeyError: 'the label [-1] is not in the [index]'
```

```
During handling of the above exception, another exception occurred:
```

```
Traceback (most recent call last):
```

```
  File "<ipython-input-1-5c89f7ac3971>", line 2, in <module>
```

```
print(df.loc[-1])
```

```
KeyError: 'the label [-1] is not in the [index]'
```

Note that passing `-1` as the `loc` will cause an error, because it is actually looking for the row index label (row number) '`-1`', which does not exist in our example. Instead, we can use a bit of Python to calculate the number of rows and pass that value into `loc`.

[Click here to view code image](#)

```
# get the last row (correctly)

# use the first value given from shape to get the number of rows

number_of_rows = df.shape[0]

# subtract 1 from the value since we want the last index value

last_row_index = number_of_rows - 1

# now do the subset using the index of the last row

print(df.loc[last_row_index])
```

```
country      Zimbabwe
```

```
continent     Africa
```

```
year          2007
```

```
lifeExp       43.487
```

```
pop           12311143
```

```
gdpPercap    469.709
```

```
Name: 1703, dtype: object
```

Alternatively, we can use the `tail` method to return the last 1 row, instead of the default 5 .

[Click here to view code image](#)

```
# there are many ways of doing what you want
```

```
print(df.tail(n=1))
```

	country	continent	year	lifeExp	pop	gdpPercap
1703	Zimbabwe	Africa	2007	43.487	12311143	469.709298

Notice that when we used `tail()` and `loc` , the results were printed out differently. Let's look at which type is returned when we use these methods.

[Click here to view code image](#)

```
subset_loc = df.loc[0]
```

```
subset_head = df.head(n=1)
```

```
# type using loc of 1 row
```

```
print(type(subset_loc))
```

```
<class 'pandas.core.series.Series'>
```

```
# type using head of 1 row
```

```
print(type(subset_head))
```

```
| <class 'pandas.core.frame.DataFrame'>
```

At the beginning of this chapter, we mentioned that Pandas introduces two new data types into Python. Depending on which method we use and how many rows we return, Pandas will return a different object. The way an object gets printed to the screen can be an indicator of the type, but it's always best to use the `type` function to be sure. We go into more details about these objects in [Chapter 2](#).

Subsetting Multiple Rows Just as for columns, we can select multiple rows.

[Click here to view code image](#)

```
# select the first, 100th, and 1000th rows  
  
# note the double square brackets similar to the syntax used to  
  
# subset multiple columns  
  
print(df.loc[[0, 99, 999]])
```

	country	continent	year	lifeExp	pop	gdpPercap
0	Afghanistan	Asia	1952	28.801	8425333	779.445314
99	Bangladesh	Asia	1967	43.453	62821884	721.186086
999	Mongolia	Asia	1967	51.253	1149500	1226.041130

1.3.2.2 Subset Rows by Row Number: iloc

`iloc` does the same thing as `loc` but is used to subset by the row index number. In our current example, `iloc` and `loc` will behave exactly the same way since the index labels are the row numbers. However, keep in mind that the index labels do not necessarily have to be row numbers.

[Click here to view code image](#)

```
# get the 2nd row
```

```
print(df.iloc[1])
```

```
country      Afghanistan
continent        Asia
year            1957
lifeExp         30.332
pop             9240934
gdpPercap       820.853
Name: 1, dtype: object
```

```
## get the 100th row
```

```
print(df.iloc[99])
```

```
country      Bangladesh
continent        Asia
year            1967
lifeExp         43.453
```

```
pop           62821884
gdpPercap     721.186
Name: 99, dtype: object
```

Note that when we put `1` into the list, we actually get the second row, rather than the first row. This follows Python's zero-indexed behavior, meaning that the first item of a container is index 0 (i.e., 0th item of the container). More details about this kind of behavior are found in [Appendices I, L, and P.](#)

With `iloc`, we can pass in the `-1` to get the last row—something we couldn't do with `loc`.

[Click here to view code image](#)

```
# using -1 to get the last row
print(df.iloc[-1])
```

```
country      Zimbabwe
continent    Africa
year         2007
lifeExp       43.487
pop          12311143
gdpPercap    469.709
Name: 1703, dtype: object
```

Just as before, we can pass in a list of integers to get multiple rows.

[Click here to view code image](#)

```
## get the first, 100th, and 1000th rows
```

```
print(df.iloc[[0, 99, 999]])
```

	country	continent	year	lifeExp	pop	gdpPercap
0	Afghanistan	Asia	1952	28.801	8425333	779.445314
99	Bangladesh	Asia	1967	43.453	62821884	721.186086
999	Mongolia	Asia	1967	51.253	1149500	1226.041130

1.3.2.3 Subsetting Rows With ix No Longer Works in Pandas v0.20

The `ix` attribute does not work in versions later than Pandas `v0.20`, since it can be confusing. Nevertheless, this section quickly reviews `ix` for completeness.

`ix` can be thought of as a combination of `loc` and `iloc`, as it allows us to subset by label or integer. By default, it searches for labels. If it cannot find the corresponding label, it falls back to using integer indexing. This can be the cause for a lot of confusion, which is why this feature has been taken out. The code using `ix` will look exactly like that written when using `loc` or `iloc`.

```
# first row
```

```
df.ix[0]
```

```
# 100th row
```

```
df.ix[99]
```

```
# 1st, 100th, and 1000th rows  
df.ix[[0, 99, 999]]
```

1.3.3 Mixing It Up

The `loc` and `iloc` attributes can be used to obtain subsets of columns, rows, or both. The general syntax for `loc` and `iloc` uses square brackets with a comma. The part to the left of the comma is the row values to subset; the part to the right of the comma is the column values to subset. That is, `df.loc[[rows], [columns]]` or `df.iloc[[rows], [columns]]`

1.3.3.1 Subsetting Columns

If we want to use these techniques to just subset columns, we must use Python's slicing syntax ([Appendix L](#)). We need to do this because if we are subsetting columns, we are getting all the rows for the specified column. So, we need a method to capture all the rows.

The Python slicing syntax uses a colon, `:`. If we have just a colon, the attribute refers to everything. So, if we just want to get the first column using the `loc` or `iloc` syntax, we can write something like `df.loc[:, [columns]]` to subset the column(s).

[Click here to view code image](#)

```
# subset columns with loc  
  
# note the position of the colon  
  
# it is used to select all rows  
  
subset = df.loc[:, ['year', 'pop']]  
  
print(subset.head())
```

```
year      pop  
0  1952    8425333  
1  1957    9240934  
2  1962    10267083  
3  1967    11537966  
4  1972    13079460
```

```
# subset columns with iloc  
  
# iloc will allow us to use integers  
  
# -1 will select the last column  
  
subset = df.iloc[:, [2, 4, -1]]  
  
print(subset.head())
```

```
year      pop    gdpPercap  
0  1952    8425333  779.445314  
1  1957    9240934  820.853030  
2  1962    10267083  853.100710  
3  1967    11537966  836.197138  
4  1972    13079460  739.981106
```

We will get an error if we don't specify `loc` and `iloc` correctly.

[Click here to view code image](#)

```
# subset columns with loc  
  
# but pass in integer values  
  
# this will cause an error  
  
subset = df.loc[:, [2, 4, -1]]  
  
print(subset.head())
```

```
| Traceback (most recent call last):  
  
|   File "<ipython-input-1-719bcb04e3c1>", line 2, in <module>  
  
|     subset = df.loc[:, [2, 4, -1]]  
  
| KeyError: 'None of [[2, 4, -1]] are in the [columns]'
```

```
# subset columns with iloc  
  
# but pass in index names  
  
# this will cause an error  
  
subset = df.iloc[:, ['year', 'pop']]  
  
print(subset.head())
```

```
| Traceback (most recent call last):  
  
|   File "<ipython-input-1-43f52fceab49>", line 2, in <module>  
  
|     subset = df.iloc[:, ['year', 'pop']]  
  
| TypeError: cannot perform reduce with flexible type
```

1.3.3.2 Subsetting Columns by Range

You can use the built-in `range` function to create a range of values in Python. This way you can specify beginning and end values, and Python will automatically create a range of values in between. By default, every value between the beginning and the end (inclusive left, exclusive right; see [Appendix L](#)) will be created, unless you specify a step ([Appendices L](#) and [P](#)). In Python 3, the `range` function returns a generator ([Appendix P](#)). If you are using Python 2, the `range` function returns a list ([Appendix I](#)), and the `xrange` function returns a generator.

If we look at the code given earlier ([Section 1.3.1.2](#)), we see that we subset columns using a list of integers. Since `range` returns a generator, we have to convert the generator to a list first.

Note that when `range(5)` is called, five integers are returned: 0 – 4.

[Click here to view code image](#)

```
# create a range of integers from 0 to 4 inclusive
```

```
small_range = list(range(5))
```

```
print(small_range)
```

```
[0, 1, 2, 3, 4]
```

```
# subset the dataframe with the range
```

```
subset = df.iloc[:, small_range]
```

```
print(subset.head())
```

	country	continent	year	lifeExp	pop
0	Afghanistan	Asia	1952	28.801	8425333

```
1  Afghanistan      Asia  1957  30.332  9240934
2  Afghanistan      Asia  1962  31.997  10267083
3  Afghanistan      Asia  1967  34.020  11537966
4  Afghanistan      Asia  1972  36.088  13079460
```

```
# create a range from 3 to 5 inclusive
```

```
small_range = list(range(3, 6))
```

```
print(small_range)
```

```
[3, 4, 5]
```

```
subset = df.iloc[:, small_range]
```

```
print(subset.head())
```

```
lifeExp      pop    gdpPercap
0   28.801    8425333  779.445314
1   30.332    9240934  820.853030
2   31.997    10267083  853.100710
3   34.020    11537966  836.197138
4   36.088    13079460  739.981106
```

QUESTION

What happens when you specify a range that's beyond the number of columns you have?

Again, note that the values are specified in a way such that the range is inclusive on the left, and exclusive on the right.

[Click here to view code image](#)

```
# create a range from 0 to 5 inclusive, every other integer
```

```
small_range = list(range(0, 6, 2))
```

```
subset = df.iloc[:, small_range]
```

```
print(subset.head())
```

	country	year	pop
0	Afghanistan	1952	8425333
1	Afghanistan	1957	9240934
2	Afghanistan	1962	10267083
3	Afghanistan	1967	11537966
4	Afghanistan	1972	13079460

Converting a generator to a list is a bit awkward; we can use the Python slicing syntax to fix this.

1.3.3.3 Slicing Columns

Python's slicing syntax, `:`, is similar to the `range` syntax. Instead of a function that specifies start, stop, and step values delimited by a comma, we separate the values with the colon.

If you understand what was going on with the `range` function earlier, then slicing can be seen as a shorthand means to the same thing.

While the range function can be used to create a generator and converted to a list of values, the colon syntax for slicing only has meaning when slicing and subsetting values, and has no inherent meaning on its own.

[Click here to view code image](#)

```
small_range = list(range(3))

subset = df.iloc[:, small_range]

print(subset.head())
```

	country	continent	year
0	Afghanistan	Asia	1952
1	Afghanistan	Asia	1957
2	Afghanistan	Asia	1962
3	Afghanistan	Asia	1967
4	Afghanistan	Asia	1972

slice the first 3 columns

```
subset = df.iloc[:, :3]

print(subset.head())
```

	country	continent	year
0	Afghanistan	Asia	1952
1	Afghanistan	Asia	1957
2	Afghanistan	Asia	1962

```
3  Afghanistan      Asia  1967
```

```
4  Afghanistan      Asia  1972
```

```
small_range = list(range(3, 6))
```

```
subset = df.iloc[:, small_range]
```

```
print(subset.head())
```

```
lifeExp      pop    gdpPercap
```

```
0    28.801    8425333   779.445314
```

```
1    30.332    9240934   820.853030
```

```
2    31.997    10267083   853.100710
```

```
3    34.020    11537966   836.197138
```

```
4    36.088    13079460   739.981106
```

```
# slice columns 3 to 5 inclusive
```

```
subset = df.iloc[:, 3:6]
```

```
print(subset.head())
```

```
lifeExp      pop    gdpPercap
```

```
0    28.801    8425333   779.445314
```

```
1    30.332    9240934   820.853030
```

```
2    31.997    10267083   853.100710
```

```
3    34.020    11537966   836.197138
```

```
4    36.088  13079460  739.981106
```

```
small_range = list(range(0, 6, 2))

subset = df.iloc[:, small_range]

print(subset.head())
```

```
        country   year      pop

0  Afghanistan  1952  8425333
1  Afghanistan  1957  9240934
2  Afghanistan  1962  10267083
3  Afghanistan  1967  11537966
4  Afghanistan  1972  13079460
```

```
# slice every other first 5 columns
```

```
subset = df.iloc[:, 0:6:2]
```

```
print(subset.head())
```

```
        country   year      pop

0  Afghanistan  1952  8425333
1  Afghanistan  1957  9240934
2  Afghanistan  1962  10267083
3  Afghanistan  1967  11537966
4  Afghanistan  1972  13079460
```

QUESTION

What happens if you use the slicing method with two colons, but leave a value out? For example, what is the result in each of the following cases?

- `df.iloc[:, 0:6:]`
 - `df.iloc[:, 0::2]`
 - `df.iloc[:, :6:2]`
 - `df.iloc[:, ::2]`
 - `df.iloc[:, ::]`
-

1.3.3.4 Subsetting Rows and Columns

We've been using the colon, `:`, in `loc` and `iloc` to the left of the comma. When we do so, we select all the rows in our dataframe. However, we can choose to put values to the left of the comma if we want to select specific rows along with specific columns.

```
# using loc  
  
print(df.loc[42, 'country'])
```

```
Angola
```

```
# using iloc  
  
print(df.iloc[42, 0])
```

```
Angola
```

Just make sure you don't forget the differences between `loc` and `iloc`.

[Click here to view code image](#)

```
# will cause an error

print(df.loc[42, 0])
```

```
Traceback (most recent call last):

  File "<ipython-input-1-2b69d7150b5e>", line 2, in <module>
    print(df.loc[42, 0])

TypeError: cannot do label indexing on <class
'pandas.core.indexes.base.Index'> with these indexers [0] of <class
'int'>
```

Now, look at how confusing `ix` can be. Good thing it no longer works.

[Click here to view code image](#)

```
# get the 43rd country in our data

df.ix[42, 'country']

# instead of 'country' I used the index 0

df.ix[42, 0]
```

1.3.3.5 Subsetting Multiple Rows and Columns

We can combine the row and column subsetting syntax with the multiple-row and multiple-column subsetting syntax to get various slices of our data.

[Click here to view code image](#)

```
# get the 1st, 100th, and 1000th rows  
  
# from the 1st, 4th, and 6th columns  
  
# the columns we are hoping to get are  
  
# country, LifeExp, and gdpPercap  
  
print(df.iloc[[0, 99, 999], [0, 3, 5]])
```

	country	lifeExp	gdpPercap
0	Afghanistan	28.801	779.445314
99	Bangladesh	43.453	721.186086
999	Mongolia	51.253	1226.041130

In my own work, I try to pass in the actual column names when subsetting data whenever possible. That approach makes the code more readable since you do not need to look at the column name vector to know which index is being called. Additionally, using absolute indexes can lead to problems if the column order gets changed for some reason. This is just a general rule of thumb, as there will be exceptions where using the index position is a better option (i.e., concatenating data in [Section 4.3](#)).

[Click here to view code image](#)

```
# if we use the column names directly,
```

```
# it makes the code a bit easier to read  
  
# note now we have to use loc, instead of iloc  
  
print(df.loc[[0, 99, 999], ['country', 'lifeExp', 'gdpPercap']])
```

	country	lifeExp	gdpPercap
0	Afghanistan	28.801	779.445314
99	Bangladesh	43.453	721.186086
999	Mongolia	51.253	1226.041130

Remember, you can use the slicing syntax on the row portion of the `loc` and `iloc` attributes.

[Click here to view code image](#)

```
print(df.loc[10:13, ['country', 'lifeExp', 'gdpPercap']])
```

	country	lifeExp	gdpPercap
10	Afghanistan	42.129	726.734055
11	Afghanistan	43.828	974.580338
12	Albania	55.230	1601.056136
13	Albania	59.280	1942.284244

1.4 Grouped and Aggregated Calculations

If you've worked with other numeric libraries or languages, you know that many basic statistic calculations either come with the library or are built into the language. Let's look at our Gapminder data again.

[Click here to view code image](#)

```
print(df.head(n=10))
```

	country	continent	year	lifeExp	pop	gdpPerCap
0	Afghanistan	Asia	1952	28.801	8425333	779.445314
1	Afghanistan	Asia	1957	30.332	9240934	820.853030
2	Afghanistan	Asia	1962	31.997	10267083	853.100710
3	Afghanistan	Asia	1967	34.020	11537966	836.197138
4	Afghanistan	Asia	1972	36.088	13079460	739.981106
5	Afghanistan	Asia	1977	38.438	14880372	786.113360
6	Afghanistan	Asia	1982	39.854	12881816	978.011439
7	Afghanistan	Asia	1987	40.822	13867957	852.395945
8	Afghanistan	Asia	1992	41.674	16317921	649.341395
9	Afghanistan	Asia	1997	41.763	22227415	635.341351

There are several initial questions that we can ask ourselves:

1. For each year in our data, what was the average life expectancy? What is the average life expectancy, population, and GDP?
2. What if we stratify the data by continent and perform the same calculations?
3. How many countries are listed in each continent?

1.4.1 Grouped Means

To answer the questions just posed, we need to perform a grouped (i.e., aggregate) calculation. In other words, we need to perform a calculation, be it an average or a frequency count, but apply it to each subset of a variable. Another way to think about grouped calculations is as a split-apply-combine process. We first split our data into various parts, then apply a function (or calculation) of our choosing to each of the split parts, and finally combine all the individual split calculations into a single dataframe. We accomplish grouped/aggregate computations by using the `groupby` method on dataframes.

[Click here to view code image](#)

```
# For each year in our data, what was the average life expectancy?  
  
# To answer this question,  
  
# we need to split our data into parts by year;  
  
# then we get the 'LifeExp' column and calculate the mean  
  
print(df.groupby('year')['lifeExp'].mean())
```

year	lifeExp
1952	49.057620
1957	51.507401
1962	53.609249
1967	55.678290
1972	57.647386
1977	59.570157
1982	61.533197

```
1987    63.212613  
1992    64.160338  
1997    65.014676  
2002    65.694923  
2007    67.007423  
  
Name: lifeExp, dtype: float64
```

Let's unpack the statement we used in this example. We first create a grouped object. Notice that if we printed the grouped dataframe, Pandas would return only the memory location.

[Click here to view code image](#)

```
grouped_year_df = df.groupby('year')  
  
print(type(grouped_year_df))  
  
<class 'pandas.core.groupby.DataFrameGroupBy'>  
  
print(grouped_year_df)  
  
<pandas.core.groupby.DataFrameGroupBy object at 0x7fe424583438>
```

From the grouped data, we can subset the columns of interest on which we want to perform our calculations. To our question, we need the `lifeExp` column. We can use the subsetting methods described in [Section 1.3.1.1](#).

[Click here to view code image](#)

```
grouped_year_df_lifeExp = grouped_year_df[ 'lifeExp' ]
```

```
print(type(grouped_year_df_lifeExp))
```

```
<class 'pandas.core.groupby.SeriesGroupBy'>
```

```
print(grouped_year_df_lifeExp)
```

```
<pandas.core.groupby.SeriesGroupBy object at 0x7fe423c9f208>
```

Notice that we now are given a series (because we asked for only one column) in which the contents of the series are grouped (in our example by year).

Finally, we know the `lifeExp` column is of type `float64`. An operation we can perform on a vector of numbers is to calculate the mean to get our final desired result.

[Click here to view code image](#)

```
mean_lifeExp_by_year = grouped_year_df_lifeExp.mean()
```

```
print(mean_lifeExp_by_year)
```

```
year
```

```
1952    49.057620
```

```
1957    51.507401
```

```
1962    53.609249
```

```
1967    55.678290
```

```
1972    57.647386
```

```
1977    59.570157  
1982    61.533197  
1987    63.212613  
1992    64.160338  
1997    65.014676  
2002    65.694923  
2007    67.007423  
  
Name: lifeExp, dtype: float64
```

We can perform a similar set of calculations for the population and GDP since they are of types `int64` and `float64`, respectively. But what if we want to group and stratify the data by more than one variable? And what if we want to perform the same calculation on multiple columns? We can build on the material earlier in this chapter by using a list!

[Click here to view code image](#)

```
# the backslash allows us to break up 1 long line of Python code  
# into multiple lines  
  
# df.groupby(['year', 'continent'])[['LifeExp', 'gdpPercap']].mean()  
  
# is the same as the following code  
  
multi_group_var = df.\  
    groupby(['year', 'continent'])\.  
    [['lifeExp', 'gdpPercap']].\  
    mean()
```

```
print(multi_group_var)
```

		lifeExp	gdpPercap
year	continent		
1952	Africa	39.135500	1252.572466
	Americas	53.279840	4079.062552
	Asia	46.314394	5195.484004
	Europe	64.408500	5661.057435
	Oceania	69.255000	10298.085650
1957	Africa	41.266346	1385.236062
	Americas	55.960280	4616.043733
	Asia	49.318544	5787.732940
	Europe	66.703067	6963.012816
	Oceania	70.295000	11598.522455
1962	Africa	43.319442	1598.078825
	Americas	58.398760	4901.541870
	Asia	51.563223	5729.369625
	Europe	68.539233	8365.486814
	Oceania	71.085000	12696.452430
1967	Africa	45.334538	2050.363801
	Americas	60.410920	5668.253496

	Asia	54.663640	5971.173374
	Europe	69.737600	10143.823757
	Oceania	71.310000	14495.021790
1972	Africa	47.450942	2339.615674
	Americas	62.394920	6491.334139
	Asia	57.319269	8187.468699
	Europe	70.775033	12479.575246
	Oceania	71.910000	16417.333380
1977	Africa	49.580423	2585.938508
	Americas	64.391560	7352.007126
	Asia	59.610556	7791.314020
	Europe	71.937767	14283.979110
	Oceania	72.855000	17283.957605
1982	Africa	51.592865	2481.592960
	Americas	66.228840	7506.737088
	Asia	62.617939	7434.135157
	Europe	72.806400	15617.896551
	Oceania	74.290000	18554.709840
1987	Africa	53.344788	2282.668991
	Americas	68.090720	7793.400261
	Asia	64.851182	7608.226508

	Europe	73.642167	17214.310727
	Oceania	75.320000	20448.040160
1992	Africa	53.629577	2281.810333
	Americas	69.568360	8044.934406
	Asia	66.537212	8639.690248
	Europe	74.440100	17061.568084
	Oceania	76.945000	20894.045885
1997	Africa	53.598269	2378.759555
	Americas	71.150480	8889.300863
	Asia	68.020515	9834.093295
	Europe	75.505167	19076.781802
	Oceania	78.190000	24024.175170
2002	Africa	53.325231	2599.385159
	Americas	72.422040	9287.677107
	Asia	69.233879	10174.090397
	Europe	76.700600	21711.732422
	Oceania	79.740000	26938.778040
2007	Africa	54.806038	3089.032605
	Americas	73.608120	11003.031625
	Asia	70.728485	12473.026870
	Europe	77.648600	25054.481636

```
Oceania    80.719500 29810.188275
```

The output data is grouped by year and continent. For each year–continent pair, we calculated the average life expectancy and average GDP. The data is also printed out a little differently. Notice the year and continent “column names” are not on the same line as the life expectancy and GDP “column names.” There is some hierachal structure between the year and continent row indices. We’ll discuss working with these types of data in more detail in [Chapter 10](#).

If you need to “flatten” the dataframe, you can use the `reset_index` method.

[Click here to view code image](#)

```
flat = multi_group_var.reset_index()
```

```
print(flat.head(15))
```

	year	continent	lifeExp	gdpPercap
0	1952	Africa	39.135500	1252.572466
1	1952	Americas	53.279840	4079.062552
2	1952	Asia	46.314394	5195.484004
3	1952	Europe	64.408500	5661.057435
4	1952	Oceania	69.255000	10298.085650
5	1957	Africa	41.266346	1385.236062
6	1957	Americas	55.960280	4616.043733
7	1957	Asia	49.318544	5787.732940
8	1957	Europe	66.703067	6963.012816

9	1957	Oceania	70.295000	11598.522455
10	1962	Africa	43.319442	1598.078825
11	1962	Americas	58.398760	4901.541870
12	1962	Asia	51.563223	5729.369625
13	1962	Europe	68.539233	8365.486814
14	1962	Oceania	71.085000	12696.452430

QUESTION

Does the order of the list we used to group the data matter?

1.4.2 Grouped Frequency Counts

Another common data-related task is to calculate frequencies. We can use the `nunique` and `value_counts` methods, respectively, to get counts of unique values and frequency counts on a Pandas Series .

[Click here to view code image](#)

```
# use the nunique (number unique)

# to calculate the number of unique values in a series

print(df.groupby('continent')['country'].nunique())
```

continent	
Africa	52
Americas	25
Asia	33

```
Europe      30
Oceania     2
Name: country, dtype: int64
```

QUESTION

What do you get if you use `value_counts` instead of `nunique`?

1.5 Basic Plot

Visualizations are extremely important in almost every step of the data process. They help us identify trends in data when we are trying to understand and clean the data, and they help us convey our final findings. More information about visualization and plotting is described in [Chapter 3](#).

Let's look at the yearly life expectancies for the world population again.

[Click here to view code image](#)

```
global_yearly_life_expectancy = df.groupby('year')[['lifeExp']].mean()

print(global_yearly_life_expectancy)
```

```
year
1952    49.057620
1957    51.507401
1962    53.609249
1967    55.678290
1972    57.647386
```

```
1977    59.570157  
1982    61.533197  
1987    63.212613  
1992    64.160338  
1997    65.014676  
2002    65.694923  
2007    67.007423  
  
Name: lifeExp, dtype: float64
```

We can use Pandas to create some basic plots as shown in [Figure 1.1](#).

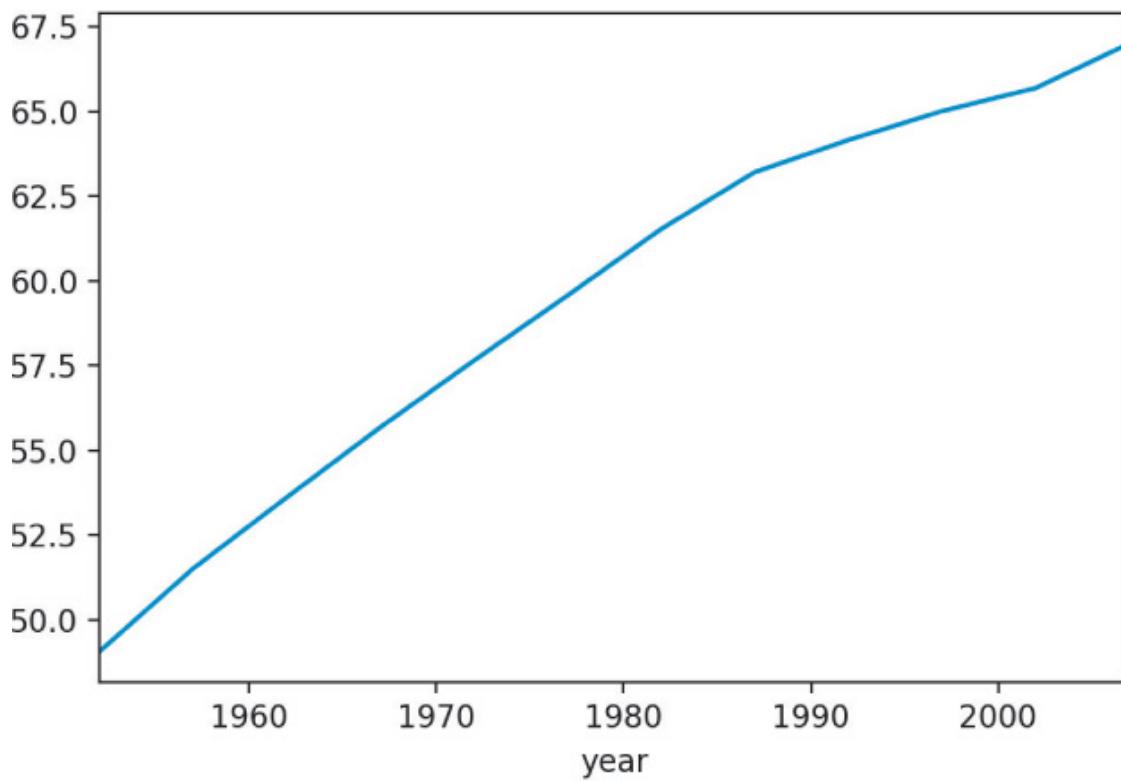


Figure 1.1 Basic plots in Pandas showing average life expectancy over time

[Click here to view code image](#)

```
global_yearly_life_expectancy.plot()
```

1.6 Conclusion

This chapter explained how to load up a simple data set and start looking at specific observations. It may seem tedious at first to look at observations this way, especially if you are already familiar with the use of a spreadsheet program. Keep in mind that when doing data analytics, the goal is to produce reproducible results, and to not repeat repetitive tasks. Scripting languages give you that ability and flexibility.

Along the way you learned about some of the fundamental programming abilities and data structures that Python has to offer. You also encountered a quick way to obtain aggregated statistics and plots. The next chapter goes into more detail about the Pandas `DataFrame` and `Series` object, as well as other ways you can subset and visualize your data.

As you work your way though this book, if there is a concept or data structure that is foreign to you, check the various appendices for more information on that topic. Many fundamental programming features of Python are covered in the appendices.

[Support](#) [Sign Out](#)

©2022 O'REILLY MEDIA, INC. [TERMS OF SERVICE](#) [PRIVACY POLICY](#)